Lab: Xv6 and Unix utilities

This lab will familiarize you with xv6 and its system calls.

Boot xv6

You can do these labs on an Athena machine or on your own computer. If you use your own computer, have a look at the **lab tools page** for setup tips.

If you use Athena, you must use an x86 machine; that is, `uname -a` should mention `i386 GNU/Linux` or `i686 GNU/Linux` or `x86_64 GNU/Linux`. You can log into a public Athena host with `ssh -X athena.dialup.mit.edu`. We have set up the appropriate compilers and simulators for you on Athena. To use them, run add -f 6.828. You must run this command every time you log in (or add it to your `~/.environment` file). If you get obscure errors while compiling or running `qemu`, check that you added the course locker.

Fetch the xv6 source for the lab and check out the `util` branch:

```
1  $ git clone git://g.csail.mit.edu/xv6-labs-2020
2  Cloning into 'xv6-labs-2020'...
3  ...
4  $ cd xv6-labs-2020
5  $ git checkout util
6  Branch 'util' set up to track remote branch 'util' from
   'origin'.
7  Switched to a new branch 'util'
```

The xv6-labs-2020 repository differs slightly from the book's xv6-riscv; it mostly adds some files. If you are curious look at the git log:

```
1  $ git log
```

The files you will need for this and subsequent lab assignments are distributed using the **Git** version control system. Above you switched to a branch (git checkout util) containing a version of xv6 tailored to this lab. To learn more about Git, take a look at the **Git user's manual**, or, you may find this **CS-oriented overview of Git**

useful. Git allows you to keep track of the changes you make to the code. For example, if you are finished with one of the exercises, and want to checkpoint your progress, you can *commit* your changes by running:

```
1  $ git commit -am 'my solution for util lab exercise 1'
2  Created commit 60d2135: my solution for util lab exercise 1
3   1 files changed, 1 insertions(+), 0 deletions(-)
4  $
```

You can keep track of your changes by using the git diff command. Running git diff will display the changes to your code since your last commit, and git diff origin/util will display the changes relative to the initial xv6-labs-2020 code. Here, `origin/xv6-labs-2020` is the name of the git branch with the initial code you downloaded for the class.

Build and run xv6:

```
1   $ make qemu
2   riscv64-unknown-elf-gcc    -c -o kernel/entry.o kernel/entry.S
3   riscv64-unknown-elf-gcc -Wall -Werror -O -fno-omit-frame-
    pointer -ggdb -DSOL_UTIL -MD -mcmodel=medany -ffreestanding -
    fno-common -nostdlib -mno-relax -I. -fno-stack-protector -fno-
    pie -no-pie   -c -o kernel/start.o kernel/start.c
4   ...
5   riscv64-unknown-elf-ld -z max-page-size=4096 -N -e main -Ttext
    0 -o user/_zombie user/zombie.o user/ulib.o user/usys.o
    user/printf.o user/umalloc.o
6   riscv64-unknown-elf-objdump -S user/_zombie > user/zombie.asm
7   riscv64-unknown-elf-objdump -t user/_zombie | sed '1,/SYMBOL
    TABLE/d; s/ .* / /; /^$/d' > user/zombie.sym
8   mkfs/mkfs fs.img README  user/xargstest.sh user/_cat user/_echo
    user/_forktest user/_grep user/_init user/_kill user/_ln
    user/_ls user/_mkdir user/_rm user/_sh user/_stressfs
    user/_usertests user/_grind user/_wc user/_zombie
9   nmeta 46 (boot, super, log blocks 30 inode blocks 13, bitmap
    blocks 1) blocks 954 total 1000
10  balloc: first 591 blocks have been allocated
11  balloc: write bitmap block at sector 45
```

```
12  qemu-system-riscv64 -machine virt -bios none -kernel
    kernel/kernel -m 128M -smp 3 -nographic -drive
    file=fs.img,if=none,format=raw,id=x0 -device virtio-blk-
    device,drive=x0,bus=virtio-mmio-bus.0
13
14  xv6 kernel is booting
15
16  hart 2 starting
17  hart 1 starting
18  init: starting sh
19  $
```

If you type `ls` at the prompt, you should see output similar to the following:

```
1   $ ls
2   .                1 1 1024
3   ..               1 1 1024
4   README           2 2 2059
5   xargstest.sh     2 3 93
6   cat              2 4 24256
7   echo             2 5 23080
8   forktest         2 6 13272
9   grep             2 7 27560
10  init             2 8 23816
11  kill             2 9 23024
12  ln               2 10 22880
13  ls               2 11 26448
14  mkdir            2 12 23176
15  rm               2 13 23160
16  sh               2 14 41976
17  stressfs         2 15 24016
18  usertests        2 16 148456
19  grind            2 17 38144
20  wc               2 18 25344
21  zombie           2 19 22408
22  console          3 20 0
```

These are the files that `mkfs` includes in the initial file system; most are programs you can run. You just ran one of them: `ls` .

xv6 has no `ps` command, but, if you type Ctrl-p, the kernel will print information about each process. If you try it now, you'll see two lines: one for `init`, and one for `sh`.

To quit qemu type: Ctrl-a x.

Grading and hand-in procedure

You can run make grade to test your solutions with the grading program. The TAs will use the same grading program to assign your lab submission a grade. Separately, we will also have check-off meetings for labs (see **Grading policy**).

The lab code comes with GNU Make rules to make submission easier. After committing your final changes to the lab, type make handin to submit your lab. For detailed instructions on how to submit see **below**.

sleep

Implement the UNIX program `sleep` for xv6; your `sleep` should pause for a user-specified number of ticks. A tick is a notion of time defined by the xv6 kernel, namely the time between two interrupts from the timer chip. Your solution should be in the file `user/sleep.c`.

Some hints:

- Before you start coding, read Chapter 1 of the **xv6 book**.
- Look at some of the other programs in `user/` (e.g., `user/echo.c`, `user/grep.c`, and `user/rm.c`) to see how you can obtain the command-line arguments passed to a program.
- If the user forgets to pass an argument, sleep should print an error message.
- The command-line argument is passed as a string; you can convert it to an integer using `atoi` (see user/ulib.c).
- Use the system call `sleep`.
- See `kernel/sysproc.c` for the xv6 kernel code that implements the `sleep` system call (look for `sys_sleep`), `user/user.h` for the C definition of `sleep`

callable from a user program, and `user/usys.S` for the assembler code that jumps from user code into the kernel for `sleep`.

- Make sure `main` calls `exit()` in order to exit your program.
- Add your `sleep` program to `UPROGS` in Makefile; once you've done that, `make qemu` will compile your program and you'll be able to run it from the xv6 shell.
- Look at Kernighan and Ritchie's book *The C programming language (second edition)* (K&R) to learn about C.

Run the program from the xv6 shell:

```
1        $ make qemu
2        ...
3        init: starting sh
4        $ sleep 10
5        (nothing happens for a little while)
6        $
7
```

Your solution is correct if your program pauses when run as shown above. Run make grade to see if you indeed pass the sleep tests.

Note that make grade runs all tests, including the ones for the assignments below. If you want to run the grade tests for one assignment, type:

```
1        $ ./grade-lab-util sleep
2
```

This will run the grade tests that match "sleep". Or, you can type:

```
1        $ make GRADEFLAGS=sleep grade
2
```

which does the same.

pingpong

Write a program that uses UNIX system calls to "ping-pong" a byte between two processes over a pair of pipes, one for each direction. The parent should send a byte to the child; the child should print ": received ping", where is its process ID, write the byte on the pipe to the parent, and exit; the parent should read the byte from the child, print ": received pong", and exit. Your solution should be in the file `user/pingpong.c` .

Some hints:

- Use `pipe` to create a pipe.
- Use `fork` to create a child.
- Use `read` to read from the pipe, and `write` to write to the pipe.
- Use `getpid` to find the process ID of the calling process.
- Add the program to `UPROGS` in Makefile.
- User programs on xv6 have a limited set of library functions available to them. You can see the list in `user/user.h` ; the source (other than for system calls) is in `user/ulib.c` , `user/printf.c` , and `user/umalloc.c` .

Run the program from the xv6 shell and it should produce the following output:

```
1    $ make qemu
2    ...
3    init: starting sh
4    $ pingpong
5    4: received ping
6    3: received pong
7    $
8
```

Your solution is correct if your program exchanges a byte between two processes and produces output as shown above.

primes

Write a concurrent version of prime sieve using pipes. This idea is due to Doug McIlroy, inventor of Unix pipes. The picture halfway down **this page** and the surrounding text explain how to do it. Your solution should be in the file `user/primes.c` .

Your goal is to use `pipe` and `fork` to set up the pipeline. The first process feeds the numbers 2 through 35 into the pipeline. For each prime number, you will arrange to create one process that reads from its left neighbor over a pipe and writes to its right neighbor over another pipe. Since xv6 has limited number of file descriptors and processes, the first process can stop at 35.

Some hints:

- Be careful to close file descriptors that a process doesn't need, because otherwise your program will run xv6 out of resources before the first process reaches 35.
- Once the first process reaches 35, it should wait until the entire pipeline terminates, including all children, grandchildren, &c. Thus the main primes process should only exit after all the output has been printed, and after all the other primes processes have exited.
- Hint: `read` returns zero when the write-side of a pipe is closed.
- It's simplest to directly write 32-bit (4-byte) `int` s to the pipes, rather than using formatted ASCII I/O.
- You should create the processes in the pipeline only as they are needed.
- Add the program to `UPROGS` in Makefile.

Your solution is correct if it implements a pipe-based sieve and produces the following output:

```
 1      $ make qemu
 2      ...
 3      init: starting sh
 4      $ primes
 5      prime 2
 6      prime 3
 7      prime 5
 8      prime 7
 9      prime 11
10      prime 13
11      prime 17
12      prime 19
13      prime 23
14      prime 29
15      prime 31
16      $
```

find

Write a simple version of the UNIX find program: find all the files in a directory tree with a specific name. Your solution should be in the file `user/find.c`.

Some hints:

- Look at user/ls.c to see how to read directories.
- Use recursion to allow find to descend into sub-directories.
- Don't recurse into "." and "..".
- Changes to the file system persist across runs of qemu; to get a clean file system run make clean and then make qemu.
- You'll need to use C strings. Have a look at K&R (the C book), for example Section 5.5.
- Note that == does not compare strings like in Python. Use strcmp() instead.
- Add the program to `UPROGS` in Makefile.

Your solution is correct if produces the following output (when the file system contains the files `b` and `a/b`):

```
1    $ make qemu
2    ...
3    init: starting sh
4    $ echo > b
5    $ mkdir a
6    $ echo > a/b
7    $ find . b
8    ./b
9    ./a/b
10   $
11
```

xargs

Write a simple version of the UNIX xargs program: read lines from the standard input and run a command for each line, supplying the line as arguments to the command. Your solution should be in the file `user/xargs.c`.

The following example illustrates xarg's behavior:

```
1    $ echo hello too | xargs echo bye
2    bye hello too
3    $
4
```

Note that the command here is "echo bye" and the additional arguments are "hello too", making the command "echo bye hello too", which outputs "bye hello too".

Please note that xargs on UNIX makes an optimization where it will feed more than argument to the command at a time. We don't expect you to make this optimization. To make xargs on UNIX behave the way we want it to for this lab, please run it with the -n option set to 1. For instance

```
1    $ echo "1\n2" | xargs -n 1 echo line
2    line 1
3    line 2
4    $
5
```

Some hints:

- Use `fork` and `exec` to invoke the command on each line of input. Use `wait` in the parent to wait for the child to complete the command.
- To read individual lines of input, read a character at a time until a newline ('\n') appears.
- kernel/param.h declares MAXARG, which may be useful if you need to declare an argv array.
- Add the program to `UPROGS` in Makefile.
- Changes to the file system persist across runs of qemu; to get a clean file system run make clean and then make qemu.

xargs, find, and grep combine well:

```
1   $ find . b | xargs grep hello
2
```

will run "grep hello" on each file named b in the directories below ".".

To test your solution for xargs, run the shell script xargstest.sh. Your solution is correct if it produces the following output:

```
1   $ make qemu
2   ...
3   init: starting sh
4   $ sh < xargstest.sh
5   $ $ $ $ $ $ hello
6   hello
7   hello
8   $ $
9
```

You may have to go back and fix bugs in your find program. The output has many `$` because the xv6 shell doesn't realize it is processing commands from a file instead of from the console, and prints a `$` for each command in the file.

Submit the lab

This completes the lab. Make sure you pass all of the make grade tests. If this lab had questions, don't forget to write up your answers to the questions in answers-*lab-name*.txt. Commit your changes (including adding answers-*lab-name*.txt) and type make handin in the lab directory to hand in your lab.Time spentCreate a new file, `time.txt`, and put in it a single integer, the number of hours you spent on the lab. Don't forget to `git add` and `git commit` the file.SubmitYou will turn in your assignments using the **submission website**. You need to request once an API key from the submission website before you can turn in any assignments or labs.

After committing your final changes to the lab, type make handin to submit your lab.

```
1   $ git commit -am "ready to submit my lab"
2   [util c2e3c8b] ready to submit my lab
3    2 files changed, 18 insertions(+), 2 deletions(-)
4
5   $ make handin
6   tar: Removing leading `/' from member names
7   Get an API key for yourself by visiting
    https://6828.scripts.mit.edu/2020/handin.py/
8   Please enter your API key: XXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXXX
9     % Total    % Received % Xferd  Average Speed   Time    Time
      Time  Current
10                                   Dload  Upload   Total   Spent
      Left  Speed
11  100 79258  100   239  100 79019   853   275k --:--:-- --:--:--
    --:--:--  276k
12  $
```

make handin will store your API key in *myapi.key*. If you need to change your API key, just remove this file and let make handin generate it again (*myapi.key* must not include newline characters).

If you run make handin and you have either uncomitted changes or untracked files, you will see output similar to the following:

```
1    M hello.c
2   ?? bar.c
3   ?? foo.pyc
4   Untracked files will not be handed in.  Continue? [y/N]
```

Inspect the above lines and make sure all files that your lab solution needs are tracked i.e. not listed in a line that begins with ??. You can cause `git` to track a new file that you create using `git add filename`.

If make handin does not work properly, try fixing the problem with the curl or Git commands. Or you can run make tarball. This will make a tar file for you, which you can then upload via our **web interface**.

- Please run `make grade` to ensure that your code passes all of the tests
- Commit any modified source code before running `make handin`
- You can inspect the status of your submission and download the submitted code at **https://6828.scripts.mit.edu/2020/handin.py/**